Lekan Adeyeri

lekanadeyeri@gmail.com | +1 (302) 433-6628 | linkedin.com/in/lekanadeyeri | hy.ms | x.com/lekanadeyeri | scholar.google.com | github.com/lekandigital

EDUCATION

Pitzer College, Cumulative GPA: 3.81/4.0; Major GPA: 3.94/4.0 — Claremont, CA
Bachelor of Arts in Computer Science and Physics — Class of 2028, on leave

Research Interests: health agents that learn from expert feedback · RL environments and verifiers built from real deployments · asynchronous, long-horizon tool use · evaluation · scientific ML

Selected Recognition: Twitter Early Bird, 1 of 28 selected from 1,000+; team placed 2nd of 5 finalists (2020) · Twitter Hackathon finalist, 60 of 3,000+ (2020) · Kleiner Perkins Engineering Fellow Finalist (2022) · Bessemer Venture Partners Fellowship Selection (2022) · Ripple Ventures Fellow (2020) · UPS Hackathon High Tech Award for worker-safety AI (2019) · Evercore Menlo Park Technology M&A Spring Selection (2024)

Selected Projects: 17 public scientific-ML experiments; distilled 1.2 MB airfoil surrogate ~5,000× faster than its tested CFD baseline at ~2.1% lift error. BodyArcade/PosePuppet: webcam-controlled 3D games (MediaPipe, Three.js) running at 111 fps with body control. Voice AI: essay on why voice assistants stalled (hy.ms/essays/2). Earlier: UPS shipping-label OCR, GE product management, Rebel One venture scouting.

EXPERIENCE

Cottonwood Technology Fund — Amsterdam, Netherlands (Hybrid)
Visiting Analyst — Jul 2026 – Present

- Only non-partner on the investment team of an ~$80M-target deep-tech fund; report directly to a partner on founder calls, sourcing and diligence. Audited 1,334 companies, identifying 961 new-to-fund leads.
- Building a Muse-style asynchronous AI analyst: the team delegates by email/WhatsApp; it plans multi-step research across papers, patents, grants and CRM, with thesis-aware memory and specialist models I'm developing.
- Typed, approval-gated tools: the model fills arguments for a fixed library of CRM actions and never writes code; each plan is approved by email, then run dry-run first, idempotent and logged.
- Long-horizon jobs run as resumable stages with retries, circuit breakers and budget gates; when capacity runs out, the agent checkpoints instead of downgrading models or relaxing validators.
- Partner corrections become calibration cases and versioned policy; a replayed past deal surfaced all four required finding classes twice. Backtesting 2,295 partner decisions showed 2,118 lacked evidence for a defensible call, shifting work to retrieval and abstention.
- Working paper: FP32/FP64 precision and optimizer stopping as the accuracy floor of PINNs and deep-learning economic solvers; implemented the study and coded a TU Delft graduate student's economics research.

One Lab — Remote
Capital Raising Associate — Jan 2026 – Jun 2026

- One of three founding-team members advising pre-seed–Series A startups; built a sandbox AI operations system where every model output cites its evidence and awaits human review; 58 offline tests.

Microsoft — Redmond, WA
New Technologists Intern, Software Engineering + Product Management — May 2021 – Aug 2021

- One of ~100 interns selected from 1,700+; built a GPT-3 chatbot prototype for Microsoft 365 on the private-beta API, months before public release (Python, Node.js); presented to M365 executives.

CAFN Lab — Stanford, CA
Computer Vision Researcher — Jun 2020 – Dec 2020

- Second author of a peer-reviewed AJNR paper on deep CNNs, which won the ASNR Cornelius G. Dyke Award, and two ISMRM abstracts; contributed cross-validation and augmentation experiments.

eResearch Technology (acquired by Clario) — Boston, MA
Software Engineering Intern — Summers 2016, 2017

- At 16, built ERT's first Alexa skill: a voice assistant so lupus patients who couldn't use touchscreens could join clinical trials by speaking; presented to the CEO and Boston Mayor's office.