Lekan Adeyeri

lekanadeyeri@gmail.com | +1 (302) 433-6628 | linkedin.com/in/lekanadeyeri | hy.ms | x.com/lekanadeyeri | scholar.google.com | github.com/lekandigital

EDUCATION

Pitzer College, Cumulative GPA: 3.81/4.0; Major GPA: 3.94/4.0 — Claremont, CA
Bachelor of Arts in Computer Science and Physics — Class of 2028, on leave

Research Interests: health agents that learn from expert feedback · RL environments and verifiers built from real deployments · asynchronous, long-horizon tool use · evaluation · scientific ML

Selected Recognition: Twitter Early Bird, 1 of 28 selected from 1,000+; team placed 2nd of 5 finalists (2020) · Twitter Hackathon finalist, 60 of 3,000+ (2020) · Kleiner Perkins Engineering Fellow Finalist (2022) · Bessemer Venture Partners Fellowship Selection (2022) · Ripple Ventures Fellow (2020) · UPS Hackathon High Tech Award for worker-safety AI (2019) · Evercore Menlo Park Technology M&A Spring Selection (2024)

Selected Projects: 17 public scientific-ML experiments; distilled 1.2 MB airfoil surrogate ~5,000× faster than its tested CFD baseline at ~2.1% lift error. BodyArcade/PosePuppet: webcam-controlled 3D games (MediaPipe, Three.js) running at 111 fps with body control. Voice AI: essay on why voice assistants stalled (hy.ms/essays/2). Earlier: UPS shipping-label OCR, GE product management, Rebel One venture scouting.

EXPERIENCE

Cottonwood Technology Fund — Amsterdam, Netherlands (Hybrid)
Visiting Analyst — Jul 2026 – Present

- Only non-partner on the investment team of an ~$80M-target deep-tech fund; report directly to a partner on founder calls, sourcing and diligence. Audited 1,334 companies, identifying 961 new-to-fund leads.
- Building a Muse-style asynchronous AI analyst: the team delegates by email/WhatsApp; it plans multi-step research across papers, patents, grants and CRM, with thesis-aware memory and specialist models I'm developing.
- Typed, approval-gated tools: the model fills arguments for a fixed library of CRM actions and never writes code; each plan is approved by email, then run dry-run first, idempotent and logged.
- Long-horizon jobs run as resumable stages with retries, circuit breakers and budget gates; when capacity runs out, the agent checkpoints instead of downgrading models or relaxing validators.
- Partner corrections become calibration cases and versioned policy; a replayed past deal surfaced all four required finding classes twice. Backtesting 2,295 partner decisions showed 2,118 lacked evidence for a defensible call, shifting work to retrieval and abstention.
- Working paper: FP32/FP64 precision and optimizer stopping as the accuracy floor of PINNs and deep-learning economic solvers; implemented the study and coded a TU Delft graduate student's economics research.

One Lab — Remote
Capital Raising Associate — Jan 2026 – Jun 2026

- One of three founding-team members advising pre-seed–Series A startups; built a sandbox AI operations system where every model output cites its evidence and awaits human review; 58 offline tests.

Microsoft Corp. — Redmond, WA
New Technologists Intern, Software Engineering + Product Management — May 2021 – Aug 2021

- One of ~100 interns selected from 1,700+; built a GPT-3 chatbot prototype for Microsoft 365 on the private-beta API, months before public release (Python, Node.js); presented to M365 executives.

CAFN Lab — Stanford, CA
Computer Vision Researcher — Jun 2020 – Dec 2020

- Second author of a peer-reviewed AJNR paper on deep CNNs, which won the ASNR Cornelius G. Dyke Award, and two ISMRM abstracts; contributed cross-validation and augmentation experiments.

eResearch Technology, Inc. (acquired by Clario) — Boston, MA
Software Engineering Intern — Summers 2016, 2017

- At 16, built ERT's first Alexa skill: a voice assistant so lupus patients who couldn't use touchscreens could join clinical trials by speaking; presented to the CEO and Boston Mayor's office.

November 5, 2026

Dear Grégoire,

Gaia2 asks what happens when the environment pushes back: an email lands mid-task, a schedule shifts, two instructions conflict. Since July, I've been building an agent that lives in that world, and I'm writing to ask you to consider me for your health, RL and agents internship at Meta Superintelligence Labs, even though I'm not a PhD student.

At Cottonwood Technology Fund, where I'm the only non-partner on the investment team, I built the fund's AI analyst, a Muse-style personal agent. The partners work with it over email and WhatsApp: it plans multi-step research across papers, patents, grants and our CRM, runs long jobs in the background, remembers each company and each partner's preferences, and acts only through typed tools a person approves. Their corrections become versioned policy and replay tests.

The most useful lesson came from evaluation. Backtesting against the partners' past decisions showed that most of those labels lacked the evidence for a defensible call. Tuning scores to them would have taught the agent to guess, so I rebuilt the work around retrieval, identity checks and explicit abstention. It's why GAIA and Gaia2 resonate with me: they measure whether an agent gets the job done in a messy environment, not whether it sounds capable.

Health is where this matters most, and it's where I started. At 16, I built a voice assistant at eResearch Technology (now Clario) so lupus patients who struggled with touchscreens could join clinical trials by speaking. At Stanford, as a software contributor, I co-authored deep-learning research on convolutional networks published in AJNR, as second author; the paper received the ASNR Cornelius G. Dyke Award. In health, a confidently wrong system is worse than none, so I care about agents that ask before acting and abstain when the evidence is thin.

I like checking whether a result comes from the method or the setup. My working paper tests whether FP32 arithmetic and optimizer stopping, not the network, set the accuracy floor of deep-learning economic solvers, extending a NeurIPS 2025 result on physics-informed networks; the study is implemented. Earlier, I built a GPT-3 prototype for Microsoft 365 on the private-beta API. What I haven't done is RL at scale, and that's what I want to learn with your team: turning the replay tests and expert corrections I've built by hand into environments and rewards for post-training.

The posting asks for a PhD student, and I'm on leave from undergraduate studies at Pitzer College. If that rules me out, I'd be grateful if you could point me to a role at MSL where this work fits. Either way, I'd love to walk you through the analyst's design and what its evaluations showed.

Thank you for reading.

Best,

Lekan Adeyeri
lekanadeyeri@gmail.com
+1 (302) 433-6628
linkedin.com/in/lekanadeyeri